HUMBERTO STEIN SHIROMOTO

BAYESIAN DATA ANALYSIS

The public is more familiar with bad design than good design. It is, in effect, conditioned to prefer bad design, because that is what it lives with. The new becomes threatening, the old reassuring.

Paul Rand, *Design, Form, and Chaos*

La perfection est atteinte, non pas lorsquíl n'y a plus rien à ajouter, mais lorsquíl n'y a plus rien à retirer.

Antoine de Saint-Exupéry

Contents

Main references:

- [Mar18]

- [**Gelman13**]

1 | Bayesian Inference

A common and useful conceptualization in statistics is to think data is generated from some true probability distribution with unknown parameters. Then, *inference* is the process of finding out the values of those parameters using just a sample (also known as a dataset) from the true probability distribution. [Mar18, pp. 12]

[**Gelman2013**]

Bayesian modelling can be summarized in three steps [Mar18]:

1. Given some data and some assumptions on how this data could have been generated, design a model by combining building blocks known as probability distributions. Most of the time these models are crude approximations, but most of the time is all we need.

2. Use Bayes' theorem to add data to the models and derive the logical consequences of combining the data and assumptions. This step is called conditioning the model on the data.

3. Criticize the model by checking whether the model makes sense according to different criteria, including the data, the expertise on the subject, and sometimes by comparing several models.

Bayesian models are also known as probabilistic models because they are built using probabilities. Why probabilities? Because probabilities are the correct mathematical tool to model uncertainty, so let's take a walk through the garden of forking paths.

Initial definitions:

$p(x)$ is called as *prior*. The prior distribution should reflect what we know about the value of the parameter before seeing the data, y.

$p(x|\theta)$ is called as *likelihood*. The likelihood is how we will introduce data in our analysis. It is an expression of the plausibility of the data given the parameters. In some texts, you will find people call this term sampling model, statistical model, or just model.

$p(\theta|x)$ is called as *posterior*. The posterior distribution is the result of the Bayesian analysis and reflects all that we know about a problem (given our data and model). The posterior is a probability distribution for the parameters in our model and not a single value. This distribution is a balance between the prior and the likelihood.

$p(\theta)$ is called as *marginal likelihood*. The last term is the marginal likelihood, also known as evidence. Formally, the marginal likelihood

is the probability of observing the data averaged over all the possible values the parameters can take (as prescribed by the prior).

Example 1.1 (Flipping a coin, [Mar18, pp. 21]). In this example, the question to be answered is: Is a coin unbiased?

To proceed, the following notation is defined

θ is the parameter.

N is the number of tosses.

y is the number of heads.

Choosing the Likelihood . Assume that only heads and tails are possible event outcomes, and that coin tosses are independent of each other. Under these assumptions, a candidate for likelihood function is the *binomial* distribution:

$$p(y|\theta, N) = \frac{N!}{y!(N-y)!}\theta^y(1-\theta)^{N-y} . \tag{1.1}$$

This is a discrete distribution returning the probability of getting y heads (or in general, successes) out of N coin tosses (or in general, trials or experiments) given a fixed value of θ.

The binomial distribution is a reasonable choice for the likelihood. One can see that θ indicates how likely it is to obtain a head when tossing a coin. Since the value of θ is unknown, it will be obtained using Bayes' theorem.

Choosing the Prior . The prior chosen is the *beta* distribution

$$p(\theta) = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)}\theta^{\alpha-1}(1-\theta)^{\beta-1} . \tag{1.2}$$

The beta distribution looks similar to the binomial except for the first term which is a normalizing constant that ensures the distribution integrates to 1, and Γ is the *gamma function*.

Why using the beta distribution as the prior?

- The beta distribution is restricted to be between 0 and 1, in the same way parameter θ is.

- The beta distribution is the conjugate prior of the binomial distribution (which we are using as the likelihood). A conjugate prior of a likelihood is a prior that, when used in combination with a given likelihood, returns a posterior with the same functional form as the prior.

Getting the Posterior. From Bayes' theorem, the posterior is proportional to the likelihood times the prior:

$$p(\theta|y) \propto p(y|\theta)p(\theta)$$
$$= \frac{N!}{y!(N-y)!}\theta^y(1-\theta)^{N-y}\frac{\Gamma(\alpha + |beta)}{\Gamma(\alpha)\Gamma(\beta)}\theta^{\alpha-1}(1-\theta)^{\beta-1} \tag{1.3}$$
$$= \theta^{y+\alpha-1}(1-\theta)^{N-y+\beta-1}$$

This expression has the same functional form of a beta distribution (except for the normalization term) with $\alpha_{posterior} = \alpha_{prior} + y$ and $\beta_{posterior} = \beta_{prior} + N - y$. In fact,

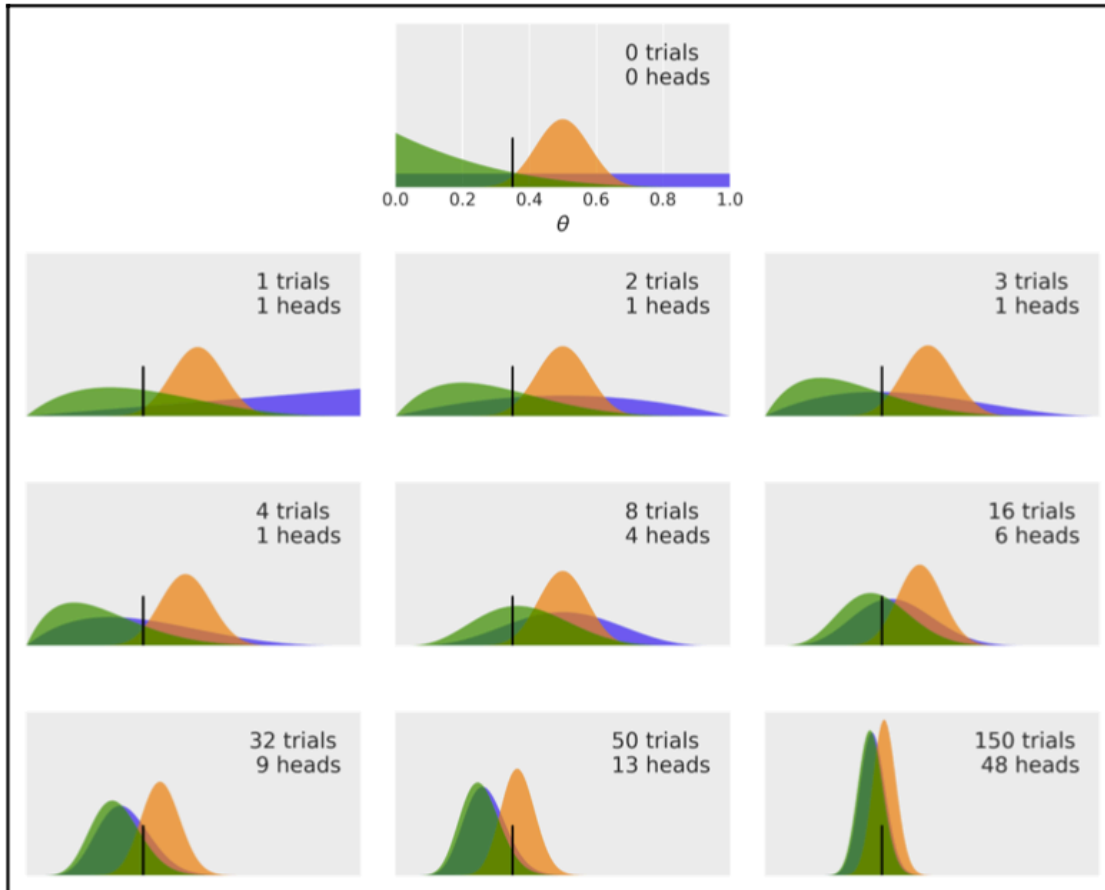$$p(\theta|y) \propto Beta(\alpha_{prior} + y, \beta_{prior} + N - y) .$$

Figure 1.1: Plot of eq. (1.3) for beta distribution parameters $(1, 1)$, $(20, 20)$, and $(1, 4)$.

Computing and Plotting the posterior. Figure 1.1 plots the distribution of the posterior for three choices of the parameters of the beta distribution: $(1, 1), (20, 20)$, and $(1, 4)$.

When, there has been 0 trials (and 0 heads), the curves plotted are the priors:

- The uniform (blue) prior represents all the possible values for the bias being equally probable a priori.

- The Gaussian-like (orange) prior is centered and concentrated around 0.5, so this prior is compatible with information indicating that the coin has more or less about the same chance of landing heads or tails.

- The skewed (green) prior puts the most weight on a tail-biased outcome.

The remaining subplots show posterior distributions for successive trials. The number of trials (or coin tosses) and the number of heads are

indicated in each subplot's legend. There is also a black vertical line at 0.35 representing the true value for θ. In practice, however, the value of θ is unknown.

Conclusion. This exercise illustrates key points of Bayesian analysis:

- The result of a Bayesian analysis is a posterior distribution, not a single value but a distribution of plausible values given the data and our model.

- The most probable value is given by the mode of the posterior (the peak of the distribution).

- The spread of the posterior is proportional to the uncertainty about the value of a parameter; the more spread out the distribution, the less certain we are.

- Intuitively, we are more confident in a result when we have observed more data supporting that result. Thus, even when numerically $\frac{1}{2} = \frac{4}{8}$, seeing four heads out of eight trials gives us more confidence that the bias is 0.5 than observing one head out of two trials. This intuition is reflected in the posterior, as you can check for yourself if you pay attention to the (blue) posterior in the third and sixth subplots; while the mode is the same, the spread (uncertainty) is larger in the third subplot than in the sixth subplot.

- Given a sufficiently large amount of data, two or more Bayesian models with different priors will tend to converge to the same result. In the limit of infinite data, no matter which prior we use, all of them will provide the same posterior. Remember that infinite is a limit and not a number, so from a practical point of view, we could get practically indistinguishably posteriors for a finite and rather small number of data points.

- How fast posteriors converge to the same distribution depends on the data and the model. In the preceding figure, we can see that the posteriors coming from the blue prior (uniform) and green prior (biased towards tails) converge faster to almost the same distribution, while it takes longer for the orange posterior (the one coming from the concentrated prior). In fact, even after 150 trials, it is somehow easy to recognize the orange posterior as a different distribution from the two others.

- Something not obvious from the figure is that we will get the same result if we update the posterior sequentially than if we do it all at once. We can compute the posterior 150 times, each time adding one more observation and using the obtained posterior as the new prior, or we can just compute one posterior for the 150 tosses at once. The result will be exactly the same. This feature not only makes perfect sense, it also leads to a natural way of updating our estimations when we get new data, a situation common in many data-analysis problems.

1.1 The Influence of Prior and How to Choose One

In general, the main problem os selectin a prior is related thinking that the prior will not allow the data "speak about itself".

Priors range in a spectrum of informative categories from non-informative to informative. Here, weakly-informative priors are considered. These are distributions that convey a small amount of information of the parameters.

Weakly-informative priors keep the posterior distributions within certain bounds. For this reason, they are also known as regularizing priors . For this reason, if a prior is reliable, there is no reason to discard it.

In particular cases, a uniformly-distributed prior is equivalent to likelihood maximization.

1.2 Probability Theory

1.3 General Probability Model

1.3.1 Choosing the Likelihood

References

[Mar18] O. Martin. *Bayesian Analysis with Python*. 2nd. Packt, 2018 (cit. on pp. 6–8).

2 | Bayesian Computation

[**Gelman:2013**]

3 | Regression Models

4 | Nonlinear Models

5 | Nonparametric Models

Contents

A | Bibliography

[CE14] is a self-contained book. It presents the concepts of linear algebra used in the book. The book starts by with linear optimisation moving to cone and semidefinite optimisation. It also contains an introduction to solving algorithms and applications to machine learning, finance, control and engineering;

[Cla13] is a more theoretical book. It contains elements of functional analysis, nonsmooth analysis and optimisation (generalised gradients). The generality of the optimisation formulation is achieved with the use of calculus of variations;

[Lib12] is a comprehensive book on the optimisation. It starts the book by introducing finite and infinite-dimensional optimisation problems. The next subject is the calculus of variations, and optimal control.

[VB96]

References

[CE14]   G. Calafiore and L. El Ghaoui. *Optimization Models*. Cambridge University Press, 2014 (cit. on p. 21).

[Cla13]   F. H. Clarke. *Functional Analysis, Calculus of Variations and Optimal Control*. Vol. 264. Graduate Texts in Mathematics. Springer, 2013 (cit. on p. 21).

[Lib12]   D. Liberzon. *Calculus of variations and optimial control: A concise introduction*. Princeton University Press, 2012 (cit. on p. 21).

[VB96]   L. Vandenberghe and S. Boyd. "Semidefinite Programming". In: *SIAM Review* 38.1 (Mar. 1996), pp. 49–95. DOI: 10 . 1137/1038003 (cit. on p. 21).

Bibliography

[CE14]   G. Calafiore and L. El Ghaoui. *Optimization Models*. Cambridge University Press, 2014 (cit. on p. 21).

[Cla13]   F. H. Clarke. *Functional Analysis, Calculus of Variations and Optimal Control*. Vol. 264. Graduate Texts in Mathematics. Springer, 2013 (cit. on p. 21).

[Lib12]   D. Liberzon. *Calculus of variations and optimial control: A concise introduction*. Princeton University Press, 2012 (cit. on p. 21).

[Mar18]   O. Martin. *Bayesian Analysis with Python*. 2nd. Packt, 2018 (cit. on pp. 6–8).

[VB96]   L. Vandenberghe and S. Boyd. "Semidefinite Programming". In: *SIAM Review* 38.1 (Mar. 1996), pp. 49–95. DOI: 10 . 1137/1038003 (cit. on p. 21).

Alphabetical Index